

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

From Logistic Regression to MLPs

Frank Hutter Abhinav Valada

University of Freiburg



Overview of Week 2

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

Recap of Logistic Regression

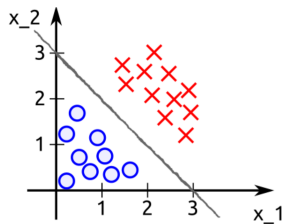
Frank Hutter Abhinav Valada

University of Freiburg



Logistic Regression: Decision Boundary

- Logistic regression is a classification method
- The decision boundary is a linear function



$$w_0 + w_1x_1 + w_2x_2 = 0$$

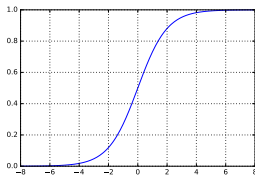
$$\mathbf{w} = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$$

Predict “ $y = 1$ ” if $\mathbf{w}^T \mathbf{x} \geq 0$ (with $x_0 = 1$ for the bias)

- Remark: with non-linear basis functions the decision boundary can also be non-linear

Logistic Regression: Probabilistic Prediction

- Logistic regression yields a probabilistic estimate
 - How likely is it that data point \mathbf{x} belongs to class 1?
- Logistic regression computes this probability as $h_{\mathbf{w}}(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x})$
 - Here, g is the logistic function $g(z) = \frac{1}{1+e^{-z}}$



- We're maximally uncertain about points on the decision boundary

$$\text{E.g., } \mathbf{w}^T \mathbf{x} = 0 \Leftrightarrow h_{\mathbf{w}}(\mathbf{x}) = 0.5$$

$$\text{E.g., } \mathbf{w}^T \mathbf{x} = 5 \Leftrightarrow h_{\mathbf{w}}(\mathbf{x}) = 0.993$$

Logistic Regression: Derivation of Cross-Entropy Loss

- The true value of y is unknown; we model it as a **random variable Y**
 - Y has a Bernoulli distribution
- Logistic regression predicts the value of Y : $h_{\mathbf{w}}(\mathbf{x}) = p_{\text{model}}(Y = 1 \mid \mathbf{x}; \mathbf{w})$
 - Model's estimated probability that $y = 1$, given that the input is \mathbf{x} and the model is parameterized by \mathbf{w}
- The actual true labels are still discrete ($y = 0$ or $y = 1$)
 - The estimated probabilities need to add to one, so:
 $p_{\text{model}}(Y = 0 \mid \mathbf{x}; \mathbf{w}) = 1 - p_{\text{model}}(Y = 1 \mid \mathbf{x}; \mathbf{w}) = 1 - h_{\mathbf{w}}(\mathbf{x})$

- Likelihood of the true data under the model:

$$p_{\text{model}}(Y = y \mid \mathbf{x}; \mathbf{w}) = \begin{cases} h_{\mathbf{w}}(\mathbf{x}) & \text{for } y = 1 \\ 1 - h_{\mathbf{w}}(\mathbf{x}) & \text{for } y = 0 \end{cases}$$

- **Cross-entropy loss**: the negative \log^1 likelihood of the true data under the model:


$$\mathcal{L}(h_{\mathbf{w}}(\mathbf{x}), y) = \begin{cases} -\log(h_{\mathbf{w}}(\mathbf{x})) & \text{for } y = 1 \\ -\log(1 - h_{\mathbf{w}}(\mathbf{x})) & \text{for } y = 0 \end{cases}$$

¹Unless mentioned otherwise, the natural logarithm is used in all formulas we will see in this lecture.


Visualization of the Cross-Entropy Loss Function

$$\mathcal{L}(h_{\mathbf{w}}(\mathbf{x}), y) = \begin{cases} -\log(h_{\mathbf{w}}(\mathbf{x})) & \text{for } y = 1 \\ -\log(1 - h_{\mathbf{w}}(\mathbf{x})) & \text{for } y = 0 \end{cases}$$

Case: $y = 1$



Case: $y = 0$



Different Way of Writing the Cross-Entropy Loss function

- Previously, we wrote the cross-entropy loss function for a single data point as:

$$\mathcal{L}(h_{\mathbf{w}}(\mathbf{x}), y) = \begin{cases} -\log(h_{\mathbf{w}}(\mathbf{x})) & \text{for } y = 1 \\ -\log(1 - h_{\mathbf{w}}(\mathbf{x})) & \text{for } y = 0 \end{cases}$$

- We can exploit that y is 0 or 1 to rewrite this in a single line:

$$\mathcal{L}(h_{\mathbf{w}}(\mathbf{x}), y) = -y \log(h_{\mathbf{w}}(\mathbf{x})) - (1 - y) \log(1 - h_{\mathbf{w}}(\mathbf{x})).$$

- Using shorthand $\hat{y} = h_{\mathbf{w}}(\mathbf{x})$ yields:

$$\mathcal{L}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

- It doesn't look like that anymore, but this is still the **negative log likelihood of the true label under the model**.

The Loss Function for the Entire Data Set

$$\mathcal{L}(\hat{y}_n, y_n) = -y_n \log \hat{y}_n - (1 - y_n) \log(1 - \hat{y}_n)$$

\rightsquigarrow Loss for a single data point $\langle \mathbf{x}_n, y_n \rangle$

- For the entire training data set $\mathcal{D}_{\text{train}} = \{\langle \mathbf{x}_1, y_1 \rangle, \dots, \langle \mathbf{x}_N, y_N \rangle\}$, the loss is:

$$L(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N y_n \log \hat{y}_n + (1 - y_n) \log(1 - \hat{y}_n)$$

- This is a **convex** function of \mathbf{w}
- We minimize it via an iterative solver (SGD)

Question to Answer for Yourself / Discuss with Friends

- Repeating a derivation: Derive the cross-entropy loss function for logistic regression:

$$\mathcal{L}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

- Application of what you just learned: What is the computational complexity (in big-O notation) of computing the cross-entropy loss

$$L(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N y_n \log \hat{y}_n + (1 - y_n) \log(1 - \hat{y}_n)$$

for logistic regression on a data set of N data points with d dimensions?

- Application of what you just learned: Which of these two predictions does cross-entropy loss prefer for a dataset $\{\langle \mathbf{x}_1, 0 \rangle, \langle \mathbf{x}_2, 1 \rangle\}$?
 - ① $\hat{y}(\mathbf{x}_1) = 1, \hat{y}(\mathbf{x}_2) = 1$, or
 - ② $\hat{y}(\mathbf{x}_1) = 0.5, \hat{y}(\mathbf{x}_2) = 0.5$?

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood**
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

Cross Entropy, KL Divergence, and Maximum Likelihood

Frank Hutter Abhinav Valada

University of Freiburg



Entropy and Cross-Entropy

- Information Entropy: $H(P) = -\mathbb{E}_{x \sim P} [\log P(x)]$
 - quantifies uncertainty about random variable X
 - for random variable with K possible outcomes:

$$H(P) = - \sum_{k=1}^K P(X = x_k) \log P(X = x_k)$$

- average number of bits² needed to code an event drawn from $P(x)$
 - Cross-Entropy: $H(P, Q) = -\mathbb{E}_{x \sim P} [\log Q(x)]$
 - for random variable with K possible outcomes:
- $$H(P, Q) = - \sum_{k=1}^K P(X = x_k) \log Q(X = x_k)$$
- average number of bits¹ needed to code an event drawn from $P(x)$ using a code that is optimized for the “wrong” distribution $Q(x)$

²bits when using \log_2 ; nats when using \ln

Kullback-Leibler (KL) Divergence

- Kullback-Leibler (KL) Divergence:

$$\begin{aligned}D_{KL}(P||Q) &= -\mathbb{E}_{\mathbf{x} \sim P} \left[\log \left(\frac{Q(x)}{P(x)} \right) \right] \\&= H(P, Q) - H(P)\end{aligned}$$

- widely used way to assess similarity of two distributions P and Q
- cross entropy minus entropy of P
- zero if $P = Q$; but not symmetric

Maximum Likelihood (ML) Estimation

- We would like to estimate a model parameter θ
- Intuitive goal: maximize the probability of some data under the model
- Consider a set of m examples $\mathbb{X} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ drawn from the (true but unknown) data-generating distribution $p_{\text{data}}(\mathbf{x})$
- Let $p_{\text{model}}(\mathbf{x}; \theta)$ be a parametric family of distributions, indexed by θ
 - We want to set θ to make this as similar to $p_{\text{data}}(\mathbf{x})$ as possible
- The maximum likelihood estimator θ_{ML} for θ is then defined as:

$$\theta_{ML} = \operatorname{argmax}_{\theta} p_{\text{model}}(\mathbb{X}; \theta) \quad (5.56)$$

Equation numbers from the very nice book “Deep Learning”

by Ian Goodfellow, Yoshua Bengio and Aaron Courville: <http://www.deeplearningbook.org/>.

Maximum Likelihood (ML) Estimation

- This maximum likelihood estimator θ_{ML} for θ also minimizes the cross entropy and the KL divergence between p_{model} and \hat{p}_{data} :

$$\theta_{ML} = \operatorname{argmax}_{\theta} p_{\text{model}}(\mathbb{X}; \theta) \quad (5.56)$$

$$= \operatorname{argmax}_{\theta} \prod_{i=1}^m p_{\text{model}}(\mathbf{x}^{(i)}; \theta) \quad (5.57)$$

$$= \operatorname{argmax}_{\theta} \sum_{i=1}^m \log p_{\text{model}}(\mathbf{x}^{(i)}; \theta) \quad (5.58)$$

$$= \operatorname{argmax}_{\theta} \mathbb{E}_{\mathbf{x} \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}; \theta) \quad (5.59)$$

$$= \operatorname{argmin}_{\theta} \mathbb{E}_{\mathbf{x} \sim \hat{p}_{\text{data}}} [-\log p_{\text{model}}(\mathbf{x}; \theta)] \quad (5.61)$$

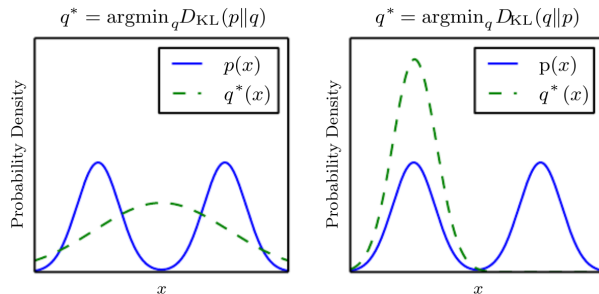
$$= \operatorname{argmin}_{\theta} H(\hat{p}_{\text{data}}, p_{\text{model}}(\cdot; \theta)) \quad (\text{cross entropy})$$

$$= \operatorname{argmin}_{\theta} \mathbb{E}_{\mathbf{x} \sim \hat{p}_{\text{data}}} [\log \hat{p}_{\text{data}}(\mathbf{x}) - \log p_{\text{model}}(\mathbf{x}; \theta)] \quad (5.60)$$

$$= \operatorname{argmin}_{\theta} D_{KL}(\hat{p}_{\text{data}}(\cdot) || p_{\text{model}}(\cdot; \theta)) \quad (\text{KL divergence})$$

Maximum Likelihood (ML) Estimation

- Interpreting the ML estimator $\theta_{ML} \in \operatorname{argmin}_{\theta} D_{KL}(\hat{p}_{\text{data}}(\cdot) || p_{\text{model}}(\cdot; \theta))$



[Goodfellow et al., 2016, p. 74, fig 3.6]

- The model parameter θ_{ML} that makes the **observed data** most likely under the model $q^* = p_{\text{model}}(\cdot; \theta_{ML})$. The left case in the figure.
- We do *not* aim for samples from q^* to be likely under p (the right case)

Question to Answer for Yourself / Discuss with Friends

- Transfer: Why would it be interesting to track the KL divergence between p_{model} and p_{data} , rather than the typically-used cross-entropy loss?

(Hint: think about how close to optimal the prediction already is.)

- Transfer: Show that the formula

$$L(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N \{y_n \log \hat{y}_n + (1 - y_n) \log(1 - \hat{y}_n)\}$$

we derived for the negative log likelihood of the Bernoulli predictions of logistic regression is just a special case of the general form

$$H(P, Q) = -\mathbb{E}_{\mathbf{x} \sim P} [\log Q(\mathbf{x})]$$

of cross-entropy.

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron**
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

Logistic Regression as a Neural Network: The Perceptron


Frank Hutter Abhinav Valada

University of Freiburg

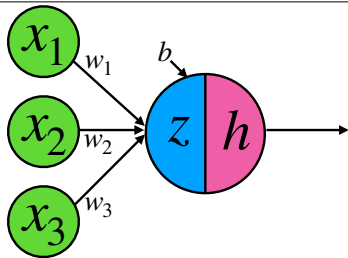


From Logistic Regression to the Perceptron

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + \exp(-(\mathbf{w}^T \mathbf{x} + w_0))} = \frac{1}{1 + \exp(-z)}$$



Logistic Regression



$$z = \mathbf{w}^T \mathbf{x} + b$$

$$h = \frac{1}{1 + \exp(-z)}$$

Perceptron

Question to Answer for Yourself / Discuss with Friends

- Repetition:
Write down the forward pass of a perceptron as a succession of two formulas.
- Repetition of previous material / transfer:
What is the computational complexity (in big-O notation) of computing the cross-entropy loss for a perceptron on a data set of N data points with d dimensions?

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons**
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

Multilayer Perceptrons

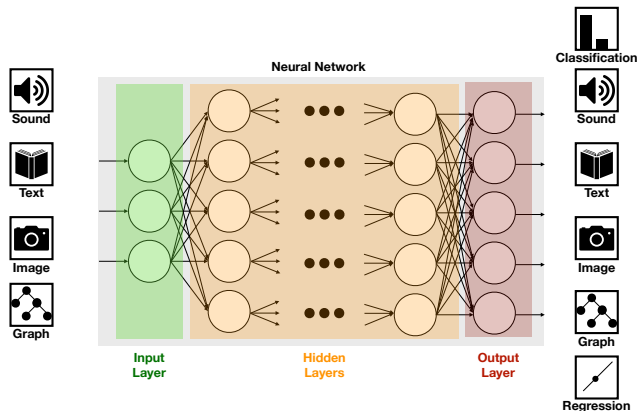
Frank Hutter Abhinav Valada

University of Freiburg

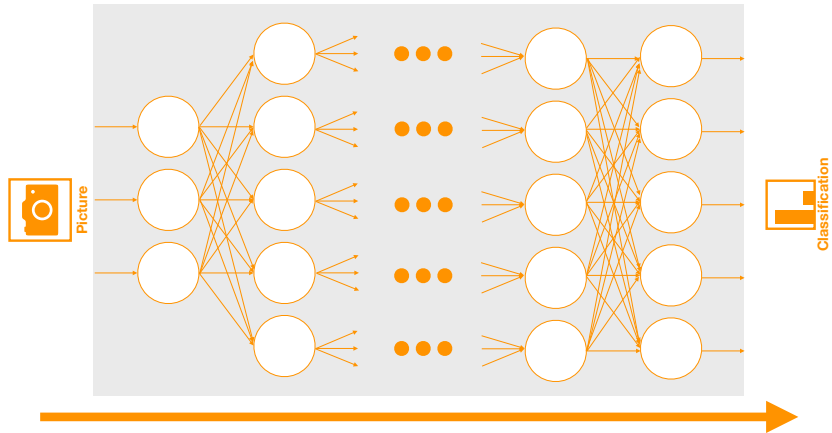


Multilayer Perceptrons (MLPs)

- We now add hidden layers
 - These layers learn nonlinear features for the final logistic regression
- Successive layers are fully-connected



Computation is Performed Layer-by-Layer



Computation in the First Hidden Layer

- For input vector \mathbf{x} , compute pre-activations $\mathbf{z}^{(1)}$ in layer 1 as

$$\mathbf{z}^{(1)} = \mathbf{W}^{(1)\top} \mathbf{x} + \mathbf{b}^{(1)}$$

- Pre-activations are transformed through a differentiable, nonlinear activation function $g^{(1)}(\cdot)$, resulting in activation vector $\mathbf{h}^{(1)}$ of the first hidden layer:

$$\mathbf{h}^{(1)} = g^{(1)}(\mathbf{z}^{(1)})$$

- The units in this layer implement the adaptable basis functions.

Computation in the Second Hidden Layer etc.

- Outputs $\mathbf{h}^{(1)}$ from layer 1 are combined linearly in the next layer 2:

$$\mathbf{z}^{(2)} = \mathbf{W}^{(2)\top} \mathbf{h}^{(1)} + \mathbf{b}^{(2)}$$

- Pre-activations $\mathbf{z}^{(2)}$ are again transformed through a nonlinear activation function $g^{(2)}$ to compute the activations $\mathbf{h}^{(2)}$:

$$\mathbf{h}^{(2)} = g^{(2)}(\mathbf{z}^{(2)})$$

- This repeats from each layer k to $k + 1$, all the way to output layer K
 - The network then outputs the output layer's activations: $\hat{\mathbf{y}} := \mathbf{h}^{(K)}$.
- E.g., for a network with one hidden layer, the overall network output $\hat{\mathbf{y}}$ for input \mathbf{x} is:

$$\hat{\mathbf{y}} = g^{(2)}(\mathbf{W}^{(2)\top} g^{(1)}(\mathbf{W}^{(1)\top} \mathbf{x} + \mathbf{b}^{(1)}) + \mathbf{b}^{(2)})$$

Question to Answer for Yourself / Discuss with Friends

- Application of what you just learned:

What is the *time* complexity (in big-O notation) of a forward pass of a single data point of input dimensionality d in an MLP with two hidden layers (of size k_1 and k_2 , respectively)?

Hint: The *time* complexity of multiplying 2 matrices with dimensions $n \times m$ and $m \times k$ is $O(nmk)$.

- Application of what you just learned:

What is the *memory* complexity (in big-O notation) of a forward pass of a single data point of input dimensionality d in an MLP with two hidden layers (of size k_1 and k_2 , respectively)?

Hint: The *memory* complexity of multiplying 2 matrices with dimensions $n \times m$ and $m \times k$ is $O(nm + mk)$.

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions**
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

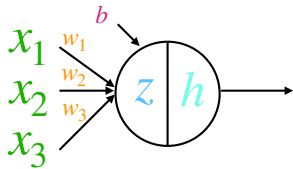
Matrix Dimensions

Frank Hutter Abhinav Valada

University of Freiburg



One Neuron, One Input Vector



$$z = w_1x_1 + w_2x_2 + w_3x_3 + b$$

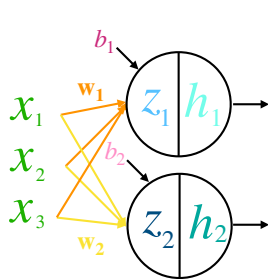
$$z = \mathbf{w}^T \mathbf{x} + b$$

A matrix representation of the equation $z = \mathbf{w}^T \mathbf{x} + b$. It shows a blue square (1x1) equals an orange row vector (1x3) multiplied by a green column vector (3x1), plus a pink square (1x1). The dimensions are labeled below each matrix: 1x1, 1x3, 3x1, and 1x1.

$$h = g(z)$$

A matrix representation of the activation function $h = g(z)$. It shows a light blue square (1x1) equals g applied to a blue square (1x1). The dimensions are labeled below each matrix: 1x1 and 1x1.

Two Neurons, One Input Vector



$$\begin{aligned}z_1 &= \mathbf{w}_1^\top \mathbf{x} + b_1 \\h_1 &= g(z_1) \\z_2 &= \mathbf{w}_2^\top \mathbf{x} + b_2 \\h_2 &= g(z_2)\end{aligned}$$



Put together into one:

$$\begin{aligned}\mathbf{z} &= \mathbf{W}^\top \mathbf{x} + \mathbf{b} \\ \mathbf{h} &= g(\mathbf{z})\end{aligned}$$

Visual illustration:

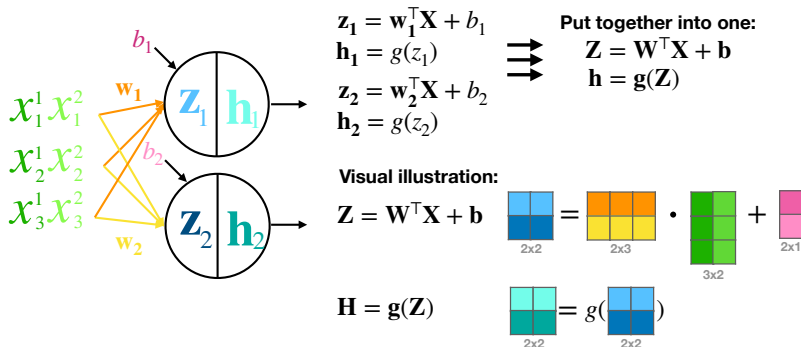
$$\mathbf{z} = \mathbf{W}^\top \mathbf{x} + \mathbf{b}$$

A visual illustration of the equation $\mathbf{z} = \mathbf{W}^\top \mathbf{x} + \mathbf{b}$. It shows a 2×1 matrix (light blue) equal to a 2×3 matrix (orange and yellow) multiplied by a 3×1 matrix (green), plus a 2×1 matrix (pink).

$$\mathbf{h} = g(\mathbf{z})$$

A visual illustration of the equation $\mathbf{h} = g(\mathbf{z})$. It shows a 2×1 matrix (cyan) equal to the function g applied to a 2×1 matrix (blue).

Two Neurons, Batch of Two Input Vectors



Two Neurons, Batch of Two Input Vectors

$$\mathbf{Z} = \mathbf{W}^T \mathbf{X} + \mathbf{b}$$

Diagram illustrating the matrix multiplication and addition for the equation $\mathbf{Z} = \mathbf{W}^T \mathbf{X} + \mathbf{b}$.

The first row shows the dimensions of the matrices:

- \mathbf{Z} (2x2) = \mathbf{W}^T (2x3) \cdot \mathbf{X} (3x2) + \mathbf{b} (2x1)

The second row shows the result of the matrix multiplication:

- \mathbf{Z} (2x2) = $\mathbf{W}^T \mathbf{X}$ (2x2) + \mathbf{b} (2x1)

The third row shows the interpretation of the matrix multiplication:

- \mathbf{Z} (2x2) = $\mathbf{W}^T \mathbf{X}$ (2x2) + \mathbf{b} (2x1) \rightarrow 2x2

The fourth row shows the final result:

- \mathbf{Z} (2x2) = $\mathbf{W}^T \mathbf{X}$ (2x2) + \mathbf{b} (2x2)

Warning: Different Common Notations in Math and in Code

- Python frameworks for Deep Learning (like **PyTorch**) use a **different notation**
 - In the slides, we follow the standard notation (e.g., in DL book) of \mathbf{x} being a column vector
 - In PyTorch, data points \mathbf{x} are row vectors
 - We will use the Pytorch notation for our **coding exercises**
- **Summary of PyTorch notation**
 - The **inputs** $\mathbf{X} \in \mathbb{R}^{N \times D}$ have N datapoints in the rows and D features in the columns
 - A single linear layer has **weight** $\mathbf{W} \in \mathbb{R}^{D \times M}$ and **bias** $\mathbf{b} \in \mathbb{R}^M$
 - The bias is **expanded** to $\mathbf{B} \in \mathbb{R}^{N \times M}$ by **repeating** it for each datapoint.
 - The **formula for output** $\mathbf{Z} \in \mathbb{R}^{N \times M}$ is then:

$$\mathbf{Z} = \mathbf{XW} + \mathbf{B}$$

Questions to Answer for Yourself / Discuss with Friends

- Repetition:

Write down the forward pass of a perceptron as a succession of two formulas, for a batch of B data points with d dimensions; for each term in the formulas, include the vector/matrix dimensions.

- Application of what you just learned:

What is the time complexity (in big-O notation) of a forward pass in an MLP with M layers of k units each, depending on the batch size B and input dimensionality d ?

- Application of what you just learned:

What is the memory complexity (in big-O notation) of a forward pass in an MLP with M layers of k units each, depending on the batch size B and input dimensionality d ?

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions**
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

Other Activation Functions and Loss Functions

Frank Hutter Abhinav Valada


University of Freiburg



Activation Functions - Examples


Logistic sigmoid activation function:

$$g_{\text{logistic}}(z) = \frac{1}{1 + \exp(-z)}$$



Logistic hyperbolic tangent activation function:

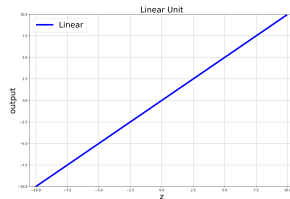
$$\begin{aligned} g_{\text{tanh}}(z) &= \tanh(z) \\ &= \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)} \end{aligned}$$



Activation Functions - Examples (cont.)

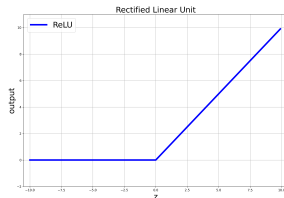
Linear activation function:

$$g_{linear}(z) = z$$



Rectified Linear (ReLU) activation function:

$$g_{relu}(z) = \max(0, z)$$




Activation Functions - Examples (cont.)

Parametric ReLU (PReLU) activation function

[He et al., 2015]:

$$PReLU(z) = \begin{cases} z, & z > 0 \\ az, & z \leq 0 \end{cases},$$

where $a > 0$ is a learnable parameter controlling the slope of the negative part




Exponential Linear Unit (ELU) activation function

[Clevert et al., 2015]:

$$ELU(z) = \begin{cases} z, & z > 0 \\ \alpha(\exp(z) - 1), & z \leq 0 \end{cases},$$

where $\alpha > 0$ controls the saturation for negative z




Activation Functions - Examples (cont.)

Gaussian Error Linear Unit (GELU) activation function [Hendrycks, Gimpel, 2016]:

$$GELU(z) = z\Phi(z),$$


where Φ is the Cumulative Distribution Function



Swish activation function [Ramachandran et al., 2017]:

$$Swish(z) = z\sigma(\beta z),$$

where $\sigma(z)$ is the sigmoid function, and $\beta \geq 0$ is a constant or trainable parameter



Output Unit Activation Functions

Depending on the task, typically:

- for regression: output neurons with linear activation
- for binary classification: output neurons with logistic/tanh activation
- for multiclass classification with K classes: use K output neurons and softmax activation

$$(\hat{\mathbf{y}}(\mathbf{x}, \mathbf{w}))_k = p(y_k = 1) = g_{softmax}(\mathbf{z})_k = \frac{\exp((\mathbf{z})_k)}{\sum_j \exp((\mathbf{z})_j)}$$

→ so for the complete output layer:

$$\hat{\mathbf{y}}(\mathbf{x}, \mathbf{w}) = \begin{bmatrix} p(y_1 = 1|\mathbf{x}) \\ p(y_2 = 1|\mathbf{x}) \\ \vdots \\ p(y_K = 1|\mathbf{x}) \end{bmatrix} = \frac{1}{\sum_{j=1}^K \exp((\mathbf{z})_j)} \exp(\mathbf{z})$$

Natural Pairing of Output Activation and Error Function

- For binary classification, use **cross-entropy error**:

$$L(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N \{y_n \log \hat{y}_n + (1 - y_n) \log(1 - \hat{y}_n)\}$$

- For linear outputs, use **mean squared error** function:

$$L(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N \{\hat{y}(\mathbf{x}_n, \mathbf{w}) - y_n\}^2$$

- For multiclass classification, use generalization of **cross-entropy error**:

$$L(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K y_{kn} \log \hat{y}_k(\mathbf{x}_n, \mathbf{w})$$

Questions to Answer for Yourself / Discuss with Friends

- Transfer: Why is a softmax function used for multiclass classification instead of simply taking the (hard)-max?
- Transfer: What would happen if you used a ReLU output activation function for regression?

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs**
- 8 Further Reading, Summary of the Week, References

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

Representational Power of MLPs

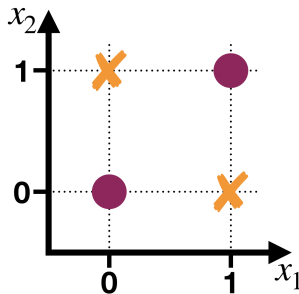
Frank Hutter Abhinav Valada

University of Freiburg



Perceptrons Can Only Model Linearly Separable Functions

- Their decision boundary is a linear function
- Famously, they cannot learn, e.g., the XOR function [Minsky, Papert, 1969]



Multilayer Perceptrons are Universal Function Approximators

Theoretical result concerning the representational power of MLPs:

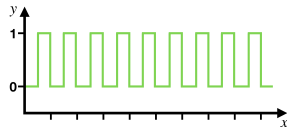
Universal Function Approximation Theorem [Cybenko, 1989]

1. Any Boolean function can be realized by an MLP with one hidden layer.
2. Any bounded continuous function can be approximated with arbitrary precision by a MLP with one hidden layer.

- The main idea of the proof:
 - Sums of (arbitrarily many) sigmoids can approximate any function
 - Similar to a Taylor expansion
- The hidden layer may have to have extremely many units
 - In the limit: infinitely many
- The theorem does not show that we can learn any function
 - It only shows that an MLP exists that approximates the function
 - It does not show that this MLP can be learned from data

The Power of Depth

- With a single hidden layer all computation has to happen in parallel
- Multiple layers allow us to re-use computation many times
 - Compositional structure of deep networks allows them to re-use pieces of computation exponentially often in terms of the network's depth.



Theorem: Depth Increases Representational Capacity Exponentially [Montufar, 2014]

A neural network with n_0 inputs and L layers of n units each, with ReLU activations can represent functions that have $\Omega((n/n_0)^{(L-1)n_0} n^{n_0})$ linear regions.

- Note: depth L is in the exponent, while width n is only in the base.

Questions to Answer for Yourself / Discuss with Friends

- Repetition: What does the universal function approximation theorem state? What does it *not* state?
- Repetition: What is the representational capacity of MLPs without a hidden layer, with one hidden layer, and with many hidden layers?

Lecture Overview

- 1 Recap of Logistic Regression
- 2 Cross Entropy, KL Divergence, and Maximum Likelihood
- 3 Logistic Regression as a Neural Network: The Perceptron
- 4 Multilayer Perceptrons
- 5 Matrix Dimensions
- 6 Other Activation Functions and Loss Functions
- 7 Representational Power of MLPs
- 8 Further Reading, Summary of the Week, References**

Foundations of Deep Learning, Winter Term 2021/22

Week 2: From Logistic Regression to MLPs

Further Reading, Summary of the Week, References

Frank Hutter Abhinav Valada

University of Freiburg



Summary by Learning Goals

- Cross-entropy loss is the negative log of the probability of predicting the correct label
- Logistic regression can be expressed by a perceptron with a logistic activation function
- Cross-entropy has a close connection to KL divergence and the maximum-likelihood estimator
- In multilayer perceptrons (MLPs) computations are performed layer-by-layer
- There are different activation functions (logistic, tanh, ReLU, linear, ELU, PReLU, etc.)
- Depending on the task, specific output unit activation functions are typically used
- Similarly, there is a natural pairing of output activations and error functions
- MLPs can represent any Boolean function, but this does not mean that they can learn any function from data
- Multiple layers (depth) increase the representational capacity of the model exponentially

Read chapter 6 of the [Deep Learning Book](#) for a detailed discussion of MLPs.

- For the latest developments on activation functions, you might find these blog posts interesting:
 - [Swish vs Mish](#)
 - [FTSwishPlus](#) and others

References

Goodfellow, I., Bengio, Y., Courville, A. (2016)

Deep Learning

M.I.T. Press.

<https://www.deeplearningbook.org/>

Minsky, M., Papert, S. (1969)

Perceptrons

M.I.T. Press

https://www.researchgate.net/publication/3081582_Review_of_'Perceptrons_An_Introduction_to_Computational_Geometry'_Minsky_M_and_Papert_S_1969

Cybenko, G. (1989)

Approximation by Superpositions of a Sigmoidal Function

Journal Mathematics of Control, Signals and Systems, 303-314

https://web.eecs.umich.edu/~cscott/smlrg/approx_by_superposition.pdf

Montúfar, G., Pascanu, R., Cho, K., Bengio, Y. (2014)

On the Number of Linear Regions of Deep Neural Networks

Conference NeurIPS 2014

<https://arxiv.org/pdf/1402.1869.pdf>

Clevert, D. A., Unterthiner, T., Hochreiter, S. (2015)

Fast and accurate deep network learning by exponential linear units (elus)

Conference ICLR 2015

<https://arxiv.org/pdf/1511.07289.pdf>

He, K., Zhang, X., Ren, S., Sun, J. (2015).

Delving deep into rectifiers: Surpassing human-level performance on imagenet classification

Conference IEEE international conference on computer vision 2015

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.725.4861&rep=rep1&type=pdf>

Hendrycks, D., Gimpel, K. (2016)

Gaussian error linear units (gelus)

<https://arxiv.org/pdf/1606.08415.pdf>

Ramachandran, P., Zoph, B., Le, Q. V. (2017)

Searching for activation functions.

<https://arxiv.org/pdf/1710.05941.pdf>