

# CS/DS 541: Class 3

Jacob Whitehill


# Hyperparameter tuning

# Hyperparameter tuning

- The values we **optimize** when training a machine learning model — e.g.,  $\mathbf{w}$  and  $b$  for linear regression — are the **parameters** of the model.
- There are also values related to the training process itself — e.g., learning rate  $\epsilon$ , batch size  $\tilde{n}$ , regularization strength  $\alpha$  — which are the **hyperparameters** of training.

# Hyperparameter tuning

- If you choose hyperparameters on the test set, you are likely deceiving yourself about how good your model is.
- **This is a subtle but very dangerous form of ML cheating.**



# Hyperparameter tuning

- Instead, you should use a separate dataset that is not part of the test set to choose hyperparameters.
- Two commonly used (and rigorous) approaches:
  - Training/validation/testing sets
  - Double cross-validation

# Training/validation/testing sets

- In an application domain with a large dataset (e.g., 100K examples), it is common to partition it into three subsets:
  - Training (typically 70-80%): optimization of parameters
  - Validation (typically 5-10%): tuning of hyperparameters
  - Testing (typically 5-10%): evaluation of the final model
- For comparison with other researchers' methods, this partition should be fixed.

# Training/validation/testing sets

- Hyperparameter tuning works as follows:
  1. For each hyperparameter configuration  $h$ :
    - Train the parameters on the **training** set using  $h$ .
    - Evaluate the model on the **validation** set.
    - If performance is better than what we got with the best  $h$  so far ( $h^*$ ), then save  $h$  as  $h^*$ .
  2. Train a model with  $h^*$ , and evaluate its accuracy  $A$  on the **testing** set. (You can train either on training data, or on training+validation data).

# Training/validation/testing sets

To what machine does the reported accuracy  $A$  correspond?

- Hyperparameter tuning works as follows:
  1. For each hyperparameter configuration  $h$ :
    - Train the parameters on the **training** set using  $h$ .
    - Evaluate the model on the **validation** set.
    - If performance is better than what we got with the best  $h$  so far ( $h^*$ ), then save  $h$  as  $h^*$ .
  2. Train a model with  $h^*$ , and evaluate its accuracy  $A$  on the **testing** set. (You can train either on training data, or on training+validation data).



# Training/validation/testing sets

To what machine does the reported accuracy  $A$  correspond?

- Hyperparameter tuning works as follows:
  1. For each hyperparameter configuration  $h$ :
    - Train the parameters on the **training** set using  $h$ .
    - Evaluate the model on the **validation** set.
    - If performance is better than what we got with the best  $h$  so far ( $h^*$ ), then save  $h$  as  $h^*$ .
  2. Train a model with  $h^*$ , and evaluate its accuracy  $A$  on the **testing** set. (You can train either on training data, or on training+validation data).

# Cross-validation

- When working with smaller datasets, cross-validation is commonly used so that we can use **all** data for training.
- Suppose we already know the best hyperparameters  $h^*$ .
- We partition the data into  $k$  “folds” of equal sizes.
- Over  $k$  iterations, we train on  $(k-1)$  folds and test on the remaining fold.
- We then iterate and compute the average accuracy over the  $k$  testing folds.

# Cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$



# Cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$



# Cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$



# Cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$



# Cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$



# Cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$





# Cross-validation

To what machine does the reported accuracy  $A$  correspond?

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$



# Cross-validation

To what machine does the reported accuracy  $A$  correspond?

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$


$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on } test$

$A = \text{Avg}[acc]$

return  $A$



None of them!

# Training/validation/testing sets


- Cross-validation does not measure the accuracy of any *single* machine.
- Instead, cross-validation gives the *expected* accuracy of a classifier that is trained on  $(k-1)/k$  of the data.

# Training/validation/testing sets

- Cross-validation does not measure the accuracy of any *single* machine.
- Instead, cross-validation gives the *expected* accuracy of a classifier that is trained on  $(k-1)/k$  of the data.
- However, we can train another model  $M$  using  $h^*$  on the entire dataset, and then report  $A$  as its accuracy.
- Since  $M$  is trained on more data than any of the cross-validation models, its *expected* accuracy should be  $\geq A$ .


# Double cross-validation

- But how do we find the best hyperparameters  $h^*$ ?
- A proper approach is to use **double cross-validation**: i.e., there will be 2 nested for-loops to iterate over “outer” and “inner” folds:




# Double cross-validation

- But how do we find the best hyperparameters  $h^*$ ?
- A proper approach is to use **double cross-validation**:, i.e., there will be 2 nested for-loops to iterate over “outer” and “inner” folds:
  - For each of the  $k$  “outer” folds, ...




# Double cross-validation

- But how do we find the best hyperparameters  $h^*$ ?
- A proper approach is to use **double cross-validation**: i.e., there will be 2 nested for-loops to iterate over “outer” and “inner” folds:
  - For each of the  $k$  “outer” folds, choose the fold-specific best  $h^*$  based on an “inner” cross-validation process.



# Double cross-validation

- But how do we find the best hyperparameters  $h^*$ ?
- A proper approach is to use **double cross-validation**: i.e., there will be 2 nested for-loops to iterate over “outer” and “inner” folds:
  - For each of the  $k$  “outer” folds, choose the fold-specific best  $h^*$  based on an “inner” cross-validation process.





# Double cross-validation

- In contrast to (single) cross-validation, it's not obvious how to finally train a model  $M$  with accuracy  $\geq A$  (i.e., the accuracy estimated by the procedure).
- One strategy: return an **ensemble** model whose output is the average of the  $k$  models' predictions...but this is rarely done.
- Double cross-validation is thus more commonly used to validate a new ML approach/architecture, rather than to build a model that is actually “delivered” to a customer.

# Dependencies among examples

- In many machine learning settings, the data are not completely independent from each other — they are *linked* in some way.
- Example:
  - Predict multiple grades for each student based on their Canvas clickstream features (# logins, # forum posts, etc.).

# Dependencies among examples

- We could partition the data into folds in different ways:
- We could randomize across all the data.
- However, if grades are correlated within each student, and if the features can reveal the students' identities, then the training data can leak information about the testing data.

Student	Major	Quiz 1	Quiz 2	Quiz 3
1	CS	45	48	42
2	Math	96	93	93
3	Chemistry	86	86	87
4	Physics	10	30	50

# Dependencies among examples

- We could partition the data into folds in different ways:
- Alternatively, we can **stratify** across students, i.e., no student appears in more than 1 fold.
- With this partition, the cross-validation accuracy estimates the model's performance on a subject *not* used for training.

Student	Major	Quiz 1	Quiz 2	Quiz 3
1	CS	45	48	42
2	Math	96	93	93
3	Chemistry	86	86	87
4	Physics	10	30	50

# Dependencies among examples

- Which strategy to use depends on how you are “marketing” your machine, e.g.:
- Our machine can predict a *new test score* for any student that the model was trained on; vs...
- Our machine can predict a test score for a *new student*.

Student	Major	Quiz 1	Quiz 2	Quiz 3
1	CS	45	48	42
2	Math	96	93	93
3	Chemistry	86	86	87
4	Physics	10	30	50

# Exercise (from d2l.ai)

- Why is the  $K$ -fold cross-validation error estimate biased?

# Linear auto-regressive (AR) models


# Linear auto-regressive (AR) models

- In some application areas (e.g., economics, finance), we have a time series of values  $x_1, x_2, \dots, x_t$ , but no “labels”  $y$ .
- Given the known values of  $x_1, x_2, \dots, x_{t-1}$ , we want to predict the value of  $x_t$ .
- A classic example is stock market price prediction.



# Linear auto-regressive (AR) models

- In one classic prediction model, we use a fixed length of history ( $p$ ) to predict the next value  $x_t$ :
  - $\hat{x}_t = W_1 x_{t-1} + W_2 x_{t-2} + \dots + W_p x_{t-p}$
- We can model (and even train) this model using the same 2-layer neural network as before:



# Auto-regression

- The essence of auto-regression is that we are using the past to predict the next future event.
- We can apply this recursively to predict infinitely into the future.
- Example for  $p=2$ , assuming we already know  $x_1, x_2$ :
  - $\hat{x}_3 = w_1 x_2 + w_2 x_1$

# Auto-regression

- The essence of auto-regression is that we are using the past to predict the next future event.
- We can apply this recursively to predict infinitely into the future.
- Example for  $p=2$ , assuming we already know  $x_1, x_2$ :
  - $\hat{x}_3 = w_1 x_2 + w_2 x_1$
  - $\hat{x}_4 = w_1 x_3 + w_2 x_2$

# Auto-regression

- The essence of auto-regression is that we are using the past to predict the next future event.
- We can apply this recursively to predict infinitely into the future.
- Example for  $p=2$ , assuming we already know  $x_1, x_2$ :
  - $\hat{x}_3 = w_1 x_2 + w_2 x_1$
  - $\hat{x}_4 = w_1 x_3 + w_2 x_2$
  - $\hat{x}_5 = w_1 x_4 + w_2 x_3$
  - ...

# Exercise

- For  $w_1=0.4$ ,  $w_2=-0.5$ ,  $x_1=0$ , and  $x_2=2$ , what are the predictions for  $x_3$ ,  $x_4$ , and  $x_5$ ?

# Exercise

- For  $w_1=0.4$ ,  $w_2=-0.5$ ,  $x_1=0$ , and  $x_2=2$ , what are the predictions for  $x_3$ ,  $x_4$ , and  $x_5$ ?
  - $\hat{x}_3 = (0.4)2 + (-0.5)0 = 0.8$
  - $\hat{x}_4 = (0.4)0.8 + (-0.5)2 = -0.68$
  - $\hat{x}_5 = (0.4)(-.68) + (-0.5)(0.8) = -0.672$

# Multivariate auto-regression

- The value  $\mathbf{x}_t$  of each time-step can also be a vector.
- In this case, each weight is a ?


# Multivariate auto-regression

- The value  $\mathbf{x}_t$  of each time-step can also be a vector.
- In this case, each weight is a matrix  $\mathbf{W}$ .
- $\hat{\mathbf{x}}_t = \mathbf{W}^{(1)} \mathbf{x}_{t-1} + \dots + \mathbf{W}^{(p)} \mathbf{x}_{t-p}$




# Multivariate auto-regression

- Suppose each observation  $\mathbf{x}_t$  has 2 components  $(x_t^a, x_t^b)$ , and that  $p=2$ .
- Here is the corresponding neural network:




# Exercise

- Recall:  $\hat{\mathbf{x}}_t = \mathbf{W}^{(1)} \mathbf{x}_{t-1} + \dots + \mathbf{W}^{(p)} \mathbf{x}_{t-p}$
- To which matrix ( $\mathbf{W}^{(1)}$ ,  $\mathbf{W}^{(2)}$ , or neither) do the first 4 edges correspond?




# Exercise

- Recall:  $\hat{\mathbf{x}}_t = \mathbf{W}^{(1)} \mathbf{x}_{t-1} + \dots + \mathbf{W}^{(p)} \mathbf{x}_{t-p}$
- To which matrix ( $\mathbf{W}^{(1)}$ ,  $\mathbf{W}^{(2)}$ , or neither) do the first 4 edges correspond?  $\mathbf{W}^{(1)}$ .



# Multivariate auto-regression

- We can alternatively represent this network with just a single matrix of weights  $\mathbf{W}$  if we “stack” the inputs:
- $\hat{\mathbf{x}}_t = \mathbf{W} [\mathbf{x}_{t-1}^\top ; \dots ; \mathbf{x}_{t-p}^\top]^\top$



# Auto-regression in deep learning

- Auto-regression is used frequently in deep learning, especially for machine translation and text generation (e.g., ChatGPT).

# Stochastic gradient descent

# Gradient descent

- With gradient descent, we only update the weights after scanning the *entire* training set.
- This is slow.
- If the training set contains 20K examples, then the gradient is an *average* over 20K images.
- How much would the gradient really change if we just used, say, 10K images? 5K images? 128 images?

$$\nabla_{\mathbf{w}} f_{\text{MSE}}(\mathbf{y}, \hat{\mathbf{y}}; \mathbf{w}) = \frac{1}{n} \mathbf{X} (\mathbf{X}^{\top} \mathbf{w} - \mathbf{y})$$

Average over entire training set.

# Stochastic gradient descent

- This is the idea behind **stochastic gradient descent** (SGD):
  - Randomly sample a small ( $\ll n$ ) **mini-batch** (or sometimes just **batch**) of training examples.
  - Estimate the gradient on just the mini-batch.
  - Update weights based on *mini-batch* gradient estimate.
  - Repeat.



# Stochastic gradient descent

- In practice, SGD is usually conducted over multiple epochs.
  - An **epoch** is a single pass through the entire training set.
- Procedure:
  1. Let  $\tilde{n} \ll n$  equal the size of the mini-batch.

# Stochastic gradient descent

- In practice, SGD is usually conducted over multiple epochs.
  - An **epoch** is a single pass through the entire training set.
- Procedure:
  1. Let  $\tilde{n} \ll n$  equal the size of the mini-batch.
  2. Randomize the order of the examples in the training set.

# Stochastic gradient descent

- In practice, SGD is usually conducted over multiple epochs.
  - An **epoch** is a single pass through the entire training set.
- Procedure:
  1. Let  $\tilde{n} \ll n$  equal the size of the mini-batch.
  2. Randomize the order of the examples in the training set.
  3. For  $e = 0$  to numEpochs:

# Stochastic gradient descent

- In practice, SGD is usually conducted over multiple epochs.
  - An **epoch** is a single pass through the entire training set.
- Procedure:
  1. Let  $\tilde{n} \ll n$  equal the size of the mini-batch.
  2. Randomize the order of the examples in the training set.
  3. For  $e = 0$  to numEpochs:
    - I. For  $i = 0$  to  $(\lceil n/\tilde{n} \rceil - 1)$  (one epoch):

# Stochastic gradient descent

- In practice, SGD is usually conducted over multiple epochs.
  - An **epoch** is a single pass through the entire training set.
- Procedure:
  1. Let  $\tilde{n} \ll n$  equal the size of the mini-batch.
  2. Randomize the order of the examples in the training set.
  3. For  $e = 0$  to numEpochs:
    - I. For  $i = 0$  to  $(\lceil n/\tilde{n} \rceil - 1)$  (one epoch):
      - A. Select a mini-batch  $\mathcal{I}$  containing the next  $\tilde{n}$  examples.

# Stochastic gradient descent

- In practice, SGD is usually conducted over multiple epochs.
  - An **epoch** is a single pass through the entire training set.
- Procedure:
  1. Let  $\tilde{n} \ll n$  equal the size of the mini-batch.
  2. Randomize the order of the examples in the training set.
  3. For  $e = 0$  to numEpochs:
    - I. For  $i = 0$  to  $(\lceil n/\tilde{n} \rceil - 1)$  (one epoch):
      - A. Select a mini-batch  $\mathcal{J}$  containing the next  $\tilde{n}$  examples.
      - B. Compute the gradient on this mini-batch:  $\frac{1}{\tilde{n}} \sum_{i \in \mathcal{J}} \nabla_{\mathbf{w}} f(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}; \mathbf{w})$

# Stochastic gradient descent

- In practice, SGD is usually conducted over multiple epochs.
  - An **epoch** is a single pass through the entire training set.
- Procedure:
  1. Let  $\tilde{n} \ll n$  equal the size of the mini-batch.
  2. Randomize the order of the examples in the training set.
  3. For  $e = 0$  to numEpochs:
    - I. For  $i = 0$  to  $(\lceil n/\tilde{n} \rceil - 1)$  (one epoch):
      - A. Select a mini-batch  $\mathcal{J}$  containing the next  $\tilde{n}$  examples.
      - B. Compute the gradient on this mini-batch:  $\frac{1}{\tilde{n}} \sum_{i \in \mathcal{J}} \nabla_{\mathbf{w}} f(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}; \mathbf{W})$
      - C. Update the weights based on the current mini-batch gradient.

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples.
- Here is how *regular* gradient descent would proceed:
  - **Initialize weights  $w^{(0)}$  to random values.**

Training  
examples

1
2
3
4
5
6
7
8



# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples.
- Here is how *regular* gradient descent would proceed:
  - Initialize weights  $\mathbf{w}^{(0)}$  to random values.
  - For each round:
    - **Compute gradient on all  $n$  examples.**

Training  
examples

1
2
3
4
5
6
7
8

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples.
- Here is how *regular* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- For each round:
  - Compute gradient on all  $n$  examples.
  - **Update weights:**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

1
2
3
4
5
6
7
8

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples.
- Here is how *regular* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- For each round:
  - **Compute gradient on all  $n$  examples.**
  - Update weights:  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

1
2
3
4
5
6
7
8

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples.
- Here is how *regular* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- For each round:
  - Compute gradient on all  $n$  examples.
  - **Update weights:**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

1
2
3
4
5
6
7
8

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:
  - **Initialize weights  $w^{(0)}$  to random values.**

Training  
examples

1
2
3
4
5
6
7
8

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:
  - Initialize weights  $\mathbf{w}^{(0)}$  to random values.
  - **Randomize the order of the training data.**

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:
  - Initialize weights  $\mathbf{w}^{(0)}$  to random values.
  - Randomize the order of the training data.
  - For each epoch ( $e=1, \dots, E$ ):  $e=1$ 
    - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
      - **Compute gradient on next  $\tilde{n}$  examples.**

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:
  - Initialize weights  $\mathbf{w}^{(0)}$  to random values.
  - Randomize the order of the training data.
  - For each epoch ( $e=1, \dots, E$ ):  **$e=1$** 
    - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
      - Compute gradient on next  $\tilde{n}$  examples.
      - **Update weights:**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2



# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- Randomize the order of the training data.
- For each epoch ( $e=1, \dots, E$ ):  $e=1$ 
  - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
    - **Compute gradient on next  $\tilde{n}$  examples.**
    - Update weights:  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:
  - Initialize weights  $\mathbf{w}^{(0)}$  to random values.
  - Randomize the order of the training data.
  - For each epoch ( $e=1, \dots, E$ ):  **$e=1$** 
    - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
      - Compute gradient on next  $\tilde{n}$  examples.
      - **Update weights:**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- Randomize the order of the training data.
- For each epoch ( $e=1, \dots, E$ ):  **$e=1$** 
  - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
    - **Compute gradient on next  $\tilde{n}$  examples.**
    - Update weights:  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:
  - Initialize weights  $\mathbf{w}^{(0)}$  to random values.
  - Randomize the order of the training data.
  - For each epoch ( $e=1, \dots, E$ ):  **$e=1$** 
    - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
      - Compute gradient on next  $\tilde{n}$  examples.
      - **Update weights:**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- Randomize the order of the training data.
- For each epoch ( $e=1, \dots, E$ ):  $e=1$ 
  - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
    - **Compute gradient on next  $\tilde{n}$  examples.**
    - Update weights:  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:
  - Initialize weights  $\mathbf{w}^{(0)}$  to random values.
  - Randomize the order of the training data.
  - For each epoch ( $e=1, \dots, E$ ):  **$e=1$** 
    - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
      - Compute gradient on next  $\tilde{n}$  examples.
      - **Update weights:**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- Randomize the order of the training data.
- For each epoch ( $e=1, \dots, E$ ):  $e=2$ 
  - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
    - **Compute gradient on next  $\tilde{n}$  examples.**
    - Update weights:  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

**Training  
examples**

4
1
3
5
7
6
8
2

# SGD versus GD: example

- Suppose our training set contains  $n=8$  examples with  $\tilde{n} = 2$ .
- Here is how *stochastic* gradient descent would proceed:

- Initialize weights  $\mathbf{w}^{(0)}$  to random values.
- Randomize the order of the training data.
- For each epoch ( $e=1, \dots, E$ ):  $e=2$ 
  - For each round ( $r=1, \dots, \lceil n/\tilde{n} \rceil$ ):
    - Compute gradient on next  $\tilde{n}$  examples.
    - Update weights:  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \epsilon \nabla_{\mathbf{w}} f$

...

Training  
examples

4
1
3
5
7
6
8
2



# Stochastic gradient descent

- Despite “noise” (statistical inaccuracy) in the mini-batch gradient estimates, we will still converge to local minimum.
- Training can be much faster than regular gradient descent because we adjust the weights *many times* per epoch.

# SGD: learning rates

- With SGD, our learning rate  $\epsilon$  needs to be **annealed** (reduced slowly over time) to guarantee convergence.
- Otherwise we might just oscillate forever in weight space.
- Necessary conditions:

$$\lim_{T \rightarrow \infty} \sum_{t=1}^T |\epsilon_t|^2 < \infty$$

Not too big: sum of squared learning rates converges.

# SGD: learning rates

- With SGD, our learning rate  $\epsilon$  needs to be **annealed** (reduced slowly over time) to guarantee convergence.
- Otherwise we might just oscillate forever in weight space.
- Necessary conditions:

$$\lim_{T \rightarrow \infty} \sum_{t=1}^T |\epsilon_t|^2 < \infty$$

$$\lim_{T \rightarrow \infty} \sum_{t=1}^T |\epsilon_t| = \infty$$

Not too small: sum of absolute learning rates grows to infinity.

# SGD: learning rates

- One common learning rate “schedule” is to multiply  $\epsilon$  by  $c \in (0, 1)$  every  $k$  rounds.
  - This is called **exponential decay**.
- Another possibility (which avoids the issue) is to set the number of epochs  $T$  to a finite number.
  - SGD may not fully converge, but the machine might still perform well.
- There are many other strategies.

# **$L_2$ Regularization**

# Regularization

- The larger the coefficients (weights)  $\mathbf{w}$  are allowed to be, the more the neural network can overfit.
- If we “encourage” the weights to be small, we can reduce overfitting.
- This is a form of **regularization** — any practice designed to improve the machine’s ability to **generalize** to new data.

# Regularization

- One of the simplest and oldest regularization techniques is to *penalize* large weights in the cost function.
- The “unregularized”  $f_{\text{MSE}}$  is:

$$f_{\text{MSE}}(\mathbf{w}) = \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

# Regularization

- One of the simplest and oldest regularization techniques is to *penalize* large weights in the cost function.

- The “unregularized”  $f_{\text{MSE}}$  is:

$$f_{\text{MSE}}(\mathbf{w}) = \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

- The  $L_2$ -regularized  $f_{\text{MSE}}$  becomes:

$$f_{\text{MSE}}(\mathbf{w}) = \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \frac{\alpha}{2n} \mathbf{w}^\top \mathbf{w}$$



# Regularization

- The closed-form solution for the optimal weight vector  $\mathbf{w}$  in  $L_2$ -regularized linear regression is then:


$$\mathbf{w} = (\mathbf{X}\mathbf{X}^\top + \alpha\mathbf{I})^{-1}\mathbf{X}\mathbf{y}$$

- Note that the inverse of the matrix in parentheses is now guaranteed to be invertible — a handy bonus.

# Optimization of ML models

# Optimization of ML models


- With linear regression, the cost function  $f_{\text{MSE}}$  has a single local minimum w.r.t. the weights  $\mathbf{w}$ :




- As long as our learning rate is small enough, we will eventually find **the optimal  $\mathbf{w}$** .

# Convex ML models

- Linear regression has a loss function that is **convex**.
- With a convex function  $f$ , *every local minimum is also a global minimum*.



convex




non-convex

- Convex functions are ideal for conducting gradient descent.

# Convexity in 1-d


- How can we tell if a 1-d function  $f$  is convex?



- What property of the slope of  $f$  ensures there is only one local minimum?

# Convexity in 1-d


- How can we tell if a 1-d function  $f$  is convex?



- What property of the slope of  $f$  ensures there is only one local minimum?
- From left to right, the slope of  $f$  *never decreases*.

# Convexity in 1-d


- How can we tell if a 1-d function  $f$  is convex?



- What property of the slope of  $f$  ensures there is only one local minimum?
- From left to right, the slope of  $f$  *never decreases*.  
==> the derivative of the slope is always non-negative.

# Convexity in 1-d

- How can we tell if a 1-d function  $f$  is convex?



- What property of the slope of  $f$  ensures there is only one local minimum?
- From left to right, the slope of  $f$  *never decreases*.  
==> the derivative of the slope is always non-negative.  
==> the second derivative of  $f$  is always non-negative.



# Convexity in higher dimensions

- For higher-dimensional  $f$ , convexity is determined by the the Hessian of  $f$ .

$$H[f] = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1 \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_m} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_m \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_m \partial x_m} \end{bmatrix}$$

- For  $f : \mathbb{R}^m \rightarrow \mathbb{R}$ ,  $f$  is convex if the Hessian matrix is positive semi-definite for *every* input  $\mathbf{x}$ .

# Positive semi-definite

- Positive semi-definite is the matrix analog of being “non-negative”.
- A real symmetric matrix **A** is **positive semi-definite (PSD)** if (equivalent conditions):

# Positive semi-definite

- Positive semi-definite is the matrix analog of being “non-negative”.
- A real symmetric matrix **A** is **positive semi-definite (PSD)** if (equivalent conditions):
  - All its eigenvalues are  $\geq 0$ .
  - In particular, if **A** happens to be diagonal, then **A** is PSD if its eigenvalues are the diagonal elements.

# Positive semi-definite

- Positive semi-definite is the matrix analog of being “non-negative”.
- A real symmetric matrix **A** is **positive semi-definite (PSD)** if (equivalent conditions):
  - All its eigenvalues are  $\geq 0$ .
    - In particular, if **A** happens to be diagonal, then **A** is PSD if its eigenvalues are the diagonal elements.
  - For every vector **v**:  $\mathbf{v}^T \mathbf{A} \mathbf{v} \geq 0$ 
    - Therefore: If there exists *any* vector **v** such that  $\mathbf{v}^T \mathbf{A} \mathbf{v} < 0$ , then **A** is *not* PSD.

# Example

- Suppose  $f(x, y) = 3x^2 + 2y^2 - 2$ .

- Then the first derivatives are:  $\frac{\partial f}{\partial x} = 6x$      $\frac{\partial f}{\partial y} = 4y$


- The Hessian matrix is therefore:

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x \partial x} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial y \partial x} & \frac{\partial^2 f}{\partial y \partial y} \end{bmatrix} = \begin{bmatrix} 6 & 0 \\ 0 & 4 \end{bmatrix}$$

- Notice that  $\mathbf{H}$  for this  $f$  does not depend on  $(x, y)$ .
- Also,  $\mathbf{H}$  is a diagonal matrix (with 6 and 4 on the diagonal). Hence, the eigenvalues are just 6 and 4. Since they are both non-negative, then  $f$  is convex.

# Example

- Graph of  $f(x, y) = 3x^2 + 2y^2 - 2$ :



# Exercises (homework)

- What about  $x^4 + xy + x^2$ ?
  - Its Hessian is:  $\mathbf{H} = \begin{bmatrix} 12x^2 + 2 & 1 \\ 1 & 0 \end{bmatrix}$
  - You either have to prove that  $\mathbf{H}$  is always PSD, or find a counter-example (homework 2).
- It can be shown (homework 2) that the MSE loss of a 2-layer linear neural network is convex.

# Convex ML models

- Prominent convex models in ML include linear regression, logistic regression, softmax regression, and support vector machines (SVM).
- However, models in deep learning are generally not convex.
- Much DL research is devoted to how to optimize the weights to deliver good generalization performance.